

Milestone 2:Data Collection & Preparation

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import f1_score
from sklearn.metrics import classification_report, confusion_matrix
import warnings
import pickle
from scipy import stats
warnings.filterwarnings('ignore')
plt.style.use('fivethirtyeight')
```

```
data=pd.read_excel("Data_Train.csv")
data.head()
```

	Airline	Date_of_Journey	Source	Destination	Route	Dep_Time	Arrival_Time	Du
0	IndiGo	24/03/2019	Banglore	New Delhi	BLR → DEL	22:20	01:10 22 Mar	
1	Air India	1/05/2019	Kolkata	Banglore	CCU → IXR → BBI → BLR → DEL	05:50	13:15	

```
category=['Airline','Source','Destination','Additional_Info']
category
```

```
['Airline', 'Source', 'Destination', 'Additional_Info']
```

```
for i in category:
    print(i,data[i].unique())
```

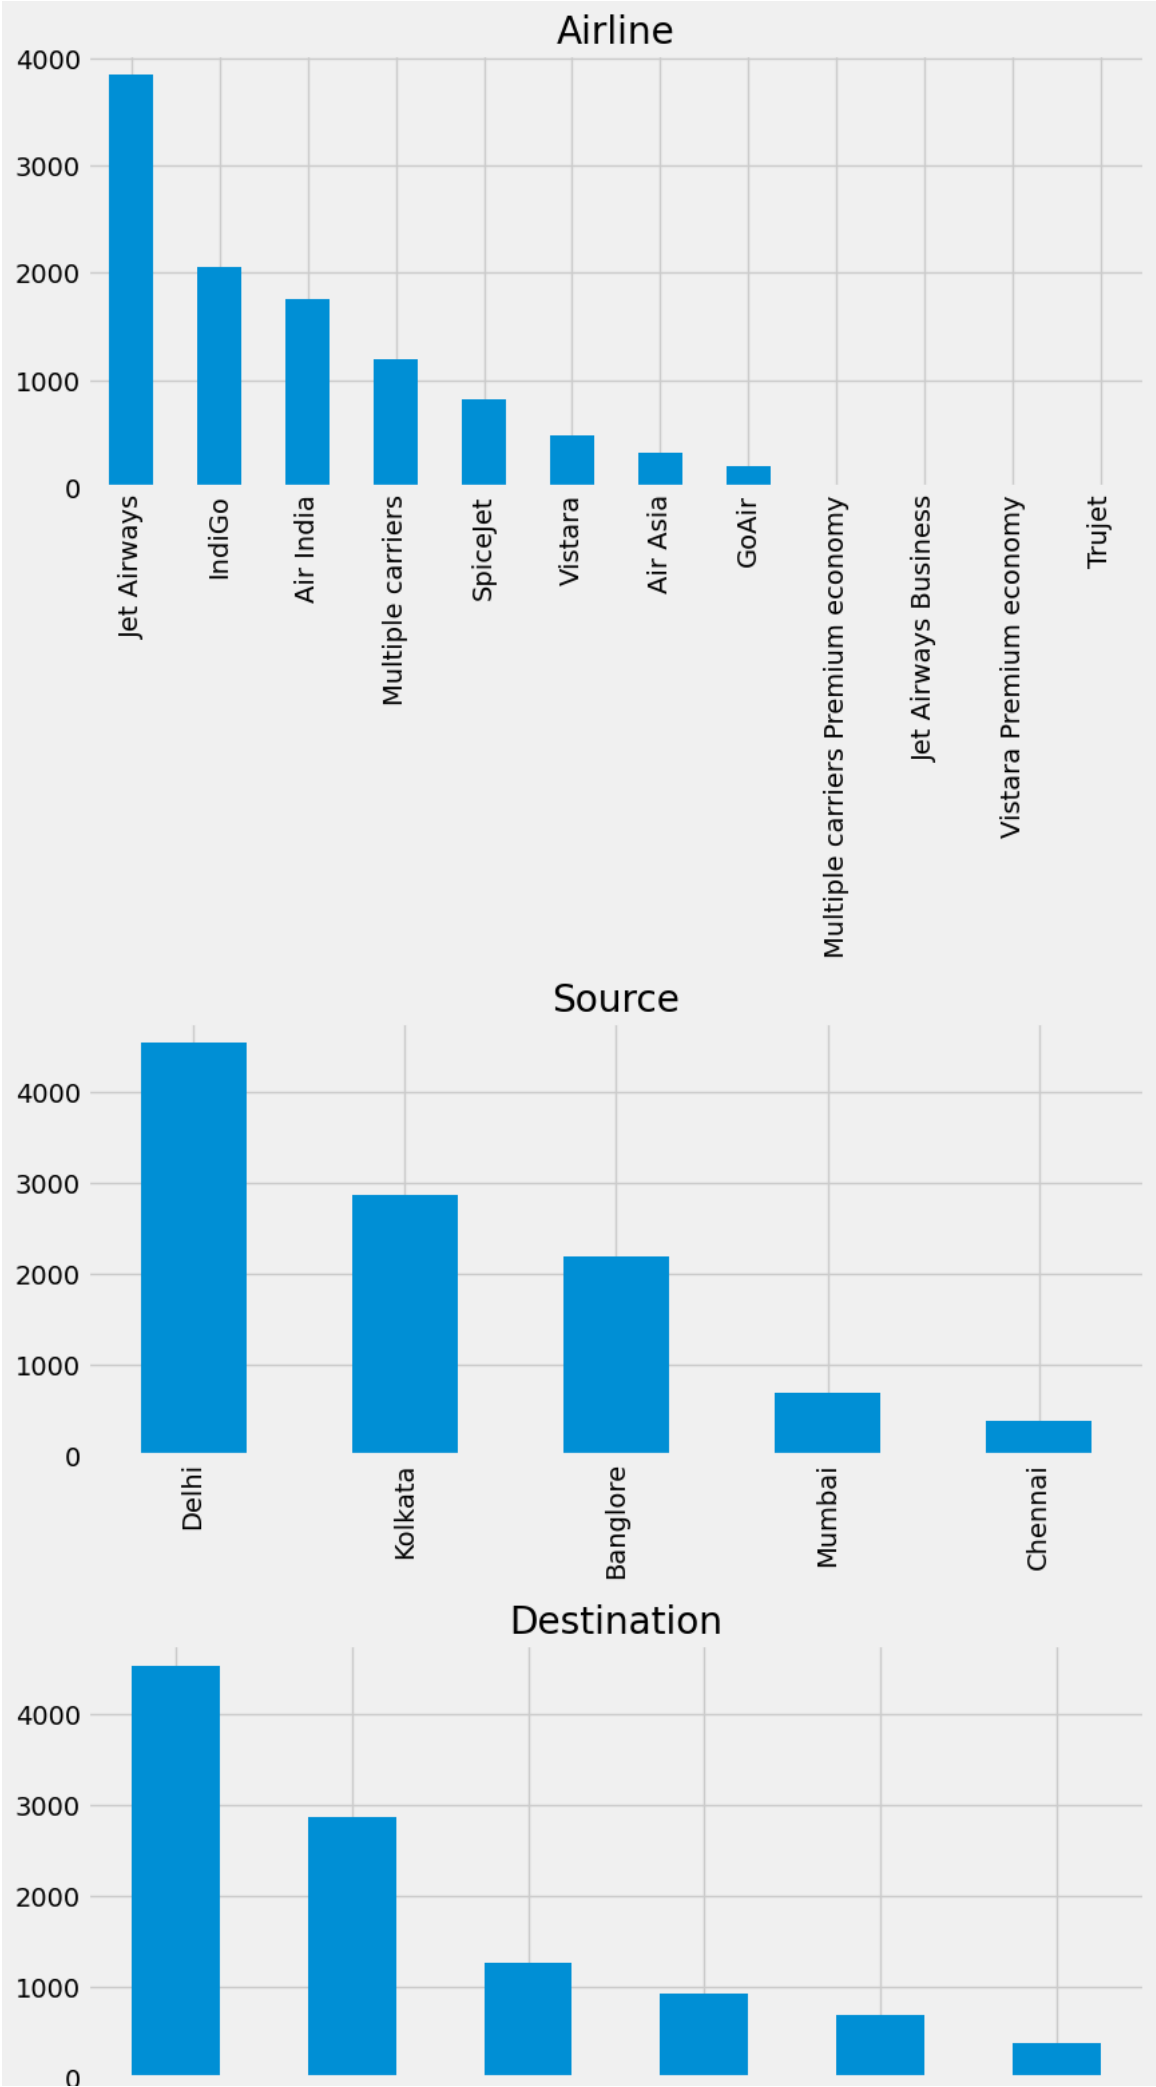
```
Airline ['IndiGo' 'Air India' 'Jet Airways' 'SpiceJet' 'Multiple carriers' 'GoAir'
'Vistara' 'Air Asia' 'Vistara Premium economy' 'Jet Airways Business'
'Multiple carriers Premium economy' 'Trujet']
Source ['Banglore' 'Kolkata' 'Delhi' 'Chennai' 'Mumbai']
Destination ['New Delhi' 'Banglore' 'Cochin' 'Kolkata' 'Delhi' 'Hyderabad']
Additional_Info ['No info' 'In-flight meal not included' 'No check-in baggage included'
'1 Short layover' 'No Info' '1 Long layover' 'Change airports'
'Business class' 'Red-eye flight' '2 Long layover']
```

```
category_cols=data.select_dtypes(include=['object']).columns
category_cols
```

```
Index(['Airline', 'Source', 'Destination', 'Route', 'Dep_Time', 'Arrival_Time',
'Duration', 'Total_Stops', 'Additional_Info', 'City1'],
      dtype='object')
```

```
#plotting a barchart for each of the categorical value
for columns in category_cols:
    plt.figure(figsize=(20,4))
```

```
plt.subplot(121)
data[columns].value_counts().plot(kind='bar')
plt.title(columns)
```



Read the Dataset

```
data.Date_of_Journey=data.Date_of_Journey.str.split('/')
data.Date_of_Journey
```

```
0      [24, 03, 2019]
1      [1, 05, 2019]
2      [9, 06, 2019]
3      [12, 05, 2019]
4      [01, 03, 2019]
...
10678   [9, 04, 2019]
10679   [27, 04, 2019]
10680   [27, 04, 2019]
10681   [01, 03, 2019]
10682   [9, 05, 2019]
Name: Date_of_Journey, Length: 10683, dtype: object
```

```
data.Total_Stops.unique()
```

```
array(['non-stop', '2 stops', '1 stop', '3 stops', nan, '4 stops'],
      dtype=object)
```

```
data.Route=data.Route.str.split('->')
data.Route
```

```
0      [BLR -> DEL]
1      [CCU -> IXR -> BBI -> BLR]
2      [DEL -> LKO -> BOM -> COK]
3      [CCU -> NAG -> BLR]
4      [BLR -> NAG -> DEL]
...
10678   [CCU -> BLR]
10679   [CCU -> BLR]
10680   [BLR -> DEL]
10681   [BLR -> DEL]
10682   [DEL -> GOI -> BOM -> COK]
Name: Route, Length: 10683, dtype: object
```

```
data['City1']=data.Route.str[0]
data['City2']=data.Route.str[1]
data['City3']=data.Route.str[2]
data['City4']=data.Route.str[3]
data['City5']=data.Route.str[4]
data['City6']=data.Route.str[5]
```

```
data.Date_of_Journey=data.Date_of_Journey.str.split('/')
data.Date_of_Journey
```

```
0      NaN
1      NaN
2      NaN
3      NaN
4      NaN
...
10678   NaN
10679   NaN
10680   NaN
10681   NaN
10682   NaN
Name: Date_of_Journey, Length: 10683, dtype: float64
```

```
data.Dep_Time=data.Dep_Time.str.split(':')
```

```
data['Dep_Time_Hour']=data.Dep_Time.str[0]
data['Dep_Time_Mins']=data.Dep_Time.str[1]
```

```
data.Arrival_Time=data.Arrival_Time.str.split('')
```

```

data['Arrival_date']=data.Arrival_Time.str[1]
data['Time_of_Arrival']=data.Arrival_Time.str[0]

100 |

data['Time_of_Arrival']=data.Time_of_Arrival.str.split(':')

data['Travel_Hours']=data.Duration.str[0]
data['Travel_Hours']=data['Travel_Hours'].str.split('h')
data['Travel_Hours']=data['Travel_Hours'].str[0]
data.Travel_Hours=data.Travel_Hours
data['Travel_Mins']=data.Duration.str[1]
data.Travel_Mins=data.Travel_Mins.str.split('m')
data.Travel_Mins=data.Travel_Mins.str[0]

```

```

data.Total_Stops.replace('non_stop',0,inplace=True)
data.Total_Stops=data.Total_Stops.str.split('')
data.Total_Stops=data.Total_Stops.str[0]

```

```

data.Total_Stops.replace('non_stop',0,inplace=True)
data.Total_Stops=data.Total_Stops.str.split('')
data.Total_stops=data.Total_Stops.str[0]

```

```
data.Additional_Info.unique()
```

```

array(['No info', 'In-flight meal not included',
       'No check-in baggage included', '1 Short layover', 'No Info',
       '1 Long layover', 'Change airports', 'Business class',
       'Red-eye flight', '2 Long layover'], dtype=object)

```

```
data.Additional_Info.replace('No Info','No Info',inplace=True)
```

Total Stops

```
data.isnull().sum()
```

```

Airline      0
Date_of_Journey  10683
Source       0
Destination  0
Route        1
Dep_Time     0
Arrival_Time 0
Duration     0
Total_Stops  1
Additional_Info 0
Price        0
City1        1
City2      10683
City3      10683
City4      10683
City5      10683
City6      10683
Dep_Time_Hour 0
Dep_Time_Mins 0
Arrival_date  0
Time_of_Arrival 0
Travel_Hours  0
Travel_Mins   0
dtype: int64

```

```
data.drop(['City4','City5','City6'],axis=1,inplace=True)
```

```

data.drop(['Date_of_Journey','Route','Dep_Time','Arrival_Time','Duration'],axis=1,inplace=True)
data.drop(['Time_of_Arrival'],axis=1,inplace=True)

```

```
data.isnull().sum()
```

```

Airline      0
Source       0
Destination  0
Total_Stops  1
Additional_Info 0
Price        0

```

```

City1          1
City2        10683
City3        10683
Dep_Time_Hour  0
Dep_Time_Mins  0
Arrival_date   0
Travel_Hours   0
Travel_Mins    0
dtype: int64

```

Replacing Missing Values

```
data['City1'].fillna('None',inplace=True)
```

```
data['Arrival_date'].fillna(data['Dep_Time_Hour'],inplace=True)
```

```
data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10683 entries, 0 to 10682
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Airline               10683 non-null  object
1   Source                10683 non-null  object
2   Destination           10683 non-null  object
3   Total_Stops           10682 non-null  object
4   Additional_Info       10683 non-null  object
5   Price                 10683 non-null  int64
6   City1                 10683 non-null  object
7   City2                 0 non-null      float64
8   City3                 0 non-null      float64
9   Dep_Time_Hour         10683 non-null  object
10  Dep_Time_Mins         10683 non-null  object
11  Arrival_date          10683 non-null  object
12  Travel_Hours          10683 non-null  object
13  Travel_Mins           10683 non-null  object
dtypes: float64(2), int64(1), object(11)
memory usage: 1.1+ MB

```

```
data['Travel_Mins'].fillna(0,inplace=True)
```

```
data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10683 entries, 0 to 10682
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Airline               10683 non-null  object
1   Source                10683 non-null  object
2   Destination           10683 non-null  object
3   Total_Stops           10682 non-null  object
4   Additional_Info       10683 non-null  object
5   Price                 10683 non-null  int64
6   Date                  10683 non-null  object
7   Month                 10683 non-null  object
8   Year                  10683 non-null  object
9   City1                 10682 non-null  object
10  City2                 10682 non-null  object
11  City3                 10683 non-null  object
12  Dep_Time_Hour         10683 non-null  object
13  Dep_Time_Mins         10683 non-null  object
14  Arrival_date          10683 non-null  object
15  Travel_Hours          10683 non-null  object
16  Travel_Mins           10683 non-null  object
dtypes: int64(1), object(16)
memory usage: 1.4+ MB

```

```

data.Dep_Time_Hour=data.Dep_Time_Hour.astype('int64')
data.Dep_Time_Mins=data.Dep_Time_Mins.astype('int64')
data.Arrival_date=data.Arrival_date.astype('int64')
data.head()

```

	Airline	Source	Destination	Total_Stops	Additional_Info	Price	City1	City2
0	IndiGo	Banglore	New Delhi	[,]	No info	3897	BLR → DEL	NaN
1	Air India	Kolkata	Banglore	[,]	No info	7662	CCU → IXR → BBI → BLR	NaN
2	Jet Airways	Delhi	Cochin	[,]	No info	13882	DEL → LKO → BOM → COK	NaN
3	IndiGo	Kolkata	Banglore	[,]	No info	6218	CCU → NAG → BLR	NaN
4	IndiGo	Banglore	New Delhi	[,]	No info	13302	BLR → NAG → DEL	NaN



```
data[data['Airline']=='Air India']
```

	Airline	Source	Destination	Total_Stops	Additional_Info	Price	City1	City2
1	Air India	Kolkata	Banglore	[.]	No info	7662	CCU → IXR → BBI → BLR	N
10	Air India	Delhi	Cochin	[.]	No info	8907	DEL → BLR → COK	N
12	Air India	Chennai	Kolkata	[.]	No info	4667	MAA → CCU	N

```
data.drop(index=6474,inplace=True,axis=0)
```

```
categorical=['Airline','Source','Destination','Additional_Info','City1']
numerical=['Total_Stops','date','Month','Year','Dep_Time_Hour','Dep_Time_Mins','Arrival_Time_Hour','Arrival_Time_Mins','Travel_Hours','Travel_Mins']
```

Label Encoding

```
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()

data.Airline=le.fit_transform(data.Airline)
data.Source=le.fit_transform(data.Source)
data.Destination=le.fit_transform(data.Destination)
```

```
data.head()
```

	Airline	Source	Destination	Total_Stops	Additional_Info	Price	City1	City2	City3	Dep_Time_Hour	Dep_Time_Mins	Arrival_date
0	3	0	5	[.]	No info	3897	BLR → DEL	NaN	NaN	22	20	
1	1	3	0	[.]	No info	7662	CCU → IXR → BBI → BLR	NaN	NaN	5	50	
2	4	2	1	[.]	No info	13882	DEL → LKO → BOM → COK	NaN	NaN	9	25	
3	3	3	0	[.]	No info	6218	CCU → NAG → BLR	NaN	NaN	18	5	
4	3	0	5	[.]	No info	13302	BLR → NAG → DEL	NaN	NaN	16	50	



Output Columns


```
data.head()
```

	Airline	Source	Destination	Total_Stops	Additional_Info	Price	City1	City2	City3	Dep_Time_Hour	Dep_Time_Mins	Arrival_date
0		3	0	5	[.]	No info	3897	BLR → DEL	NaN	NaN	22	20
1		1	3	0	[.]	No info	7662	CCU → IXR → BBI → BLR → DEL → LKO → BOM → COK → CCU → NAG → BLR → BLR → NAG → DEL	NaN	NaN	5	50
2		4	2	1	[.]	No info	13882		NaN	NaN	9	25
3		3	3	0	[.]	No info	6218		NaN	NaN	18	5
4		3	0	5	[.]	No info	13302		NaN	NaN	16	50



```
data=data[['Airline','Source','Destination','Dep_Time_Hour','Dep_Time_Mins','Arrival_date','Price']]
```

```
data.head()
```

	Airline	Source	Destination	Dep_Time_Hour	Dep_Time_Mins	Arrival_date	Price
0		3	0	5	22	20	0 3897
1		1	3	0	5	50	1 7662
2		4	2	1	9	25	0 13882
3		3	3	0	18	5	2 6218
4		3	0	5	16	50	2 13302



```
data.describe()
```

	Airline	Source	Destination	Dep_Time_Hour	Dep_Time_Mins	Arrival_date	Price
count	10682.000000	10682.000000	10682.000000	10682.000000	10682.000000	10682.000000	10682.000000
mean		3.966205	1.952069	1.435967	12.490358	24.408819	0.896836 9086.292735
std		2.352090	1.177110	1.474773	5.748819	18.767225	0.711845 4610.885695
min		0.000000	0.000000	0.000000	0.000000	0.000000	0.000000 1759.000000
25%		3.000000	2.000000	0.000000	8.000000	5.000000	0.000000 5277.000000
50%		4.000000	2.000000	1.000000	11.000000	25.000000	1.000000 8372.000000
75%		4.000000	3.000000	2.000000	18.000000	40.000000	1.000000 12373.000000
max		11.000000	4.000000	5.000000	23.000000	55.000000	2.000000 79512.000000



Milestone 3:Explortory Data Analysis

```
import seaborn as sns
c=1
plt.figure(figsize=(20,45))
```

<Figure size 2000x4500 with 0 Axes>
 <Figure size 2000x4500 with 0 Axes>

```
for i in categorical:
    plt.subplot(6,3,c)

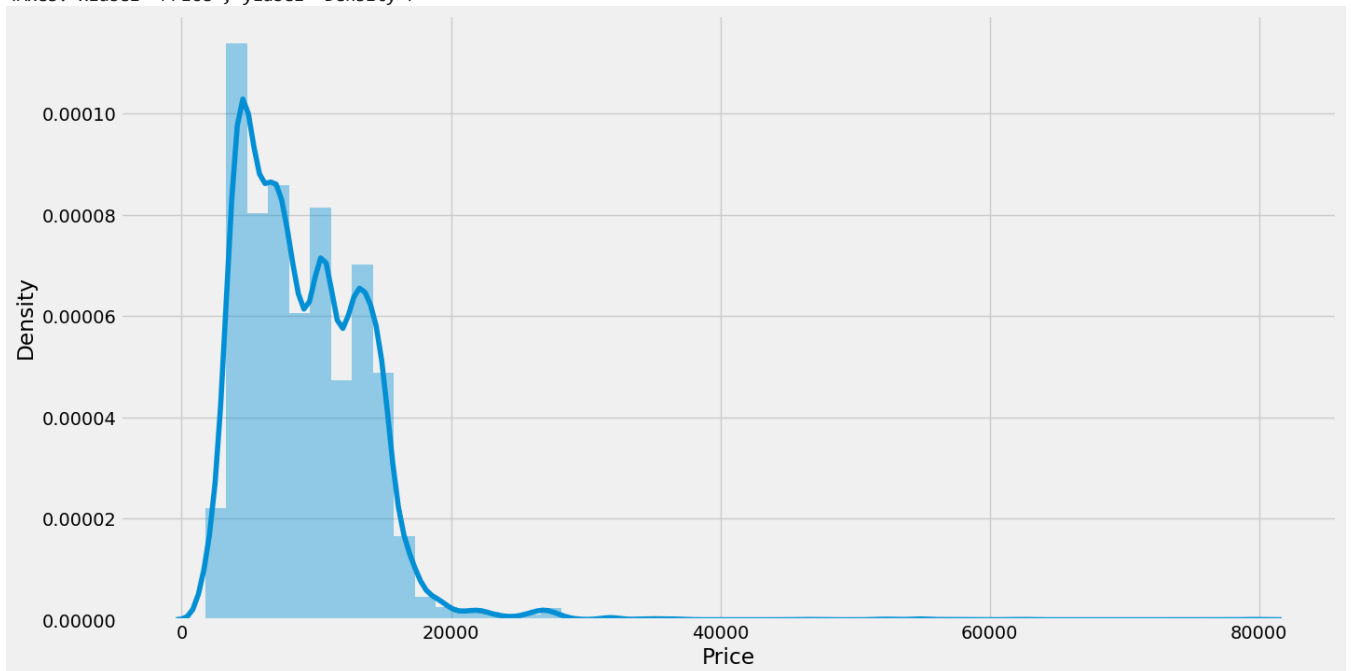
    plt.xticks(rotation=90)
    plt.tight_layout(pad=3.0)
    c=c+1
```



```
plt.show()
```


```
plt.figure(figsize=(15,8))
sns.distplot(data.Price)
```

<Axes: xlabel='Price', ylabel='Density'>




```
sns.heatmap(data.corr(),annot=True)
```

<Axes: >



```
# Detecting the Outliers
import seaborn as sns
sns.boxplot(data['Price'])
```

<Axes: >



```
y=data['Price']
x=data.drop(columns=['Price'],axis=1)
```

```
### Scaling the Data
```

```
from sklearn.preprocessing import StandardScaler
ss=StandardScaler
```

```
x_scaled=ss.fit_transform
```

```
data.head()
```

	Airline	Source	Destination	Dep_Time_Hour	Dep_Time_Mins	Arrival_date	Price
0	3	0	5	22	20	0	3897
1	1	3	0	5	50	1	7662
2	4	2	1	9	25	0	13882
3	3	3	0	18	5	2	6218
4	3	0	5	16	50	2	13302

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=42)
```

```
x_train.head()
```

	Airline	Source	Destination	Dep_Time_Hour	Dep_Time_Mins	Arrival_date	
10005	6	2	1	8	30	1	
3684	4	2	1	11	30	1	
1034	8	2	1	15	45	2	
3909	6	2	1	12	50	0	



x_train.shape

(8545, 6)

```
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor, AdaBoostRegressor
rfr=RandomForestRegressor()
gb=GradientBoostingRegressor()
ad=AdaBoostRegressor()
```

Milestone 4:Model Building

```
from sklearn.metrics import r2_score,mean_absolute_error,mean_squared_error
```

Regression Model

```
from sklearn.neighbors import KNeighborsRegressor
from sklearn.svm import SVR
from sklearn.tree import DecisionTreeRegressor
```

```
from sklearn.metrics import r2_score,mean_absolute_error,mean_squared_error
```

```
knn=KNeighborsRegressor()
svr=SVR()
dt=DecisionTreeRegressor()
```

```
for i in [knn,svr,dt]:
    i.fit(x_train,y_train)

    y_pred=i.predict(x_test)
    test_score=r2_score(y_test,y_pred)
    train_score=r2_score(y_train,i.predict(x_train))
    if abs(train_score-test_score)<=0.1:
        print(i)
        print('R2 Score is',r2_score(y_test,y_pred))
        print('R2 Score for train data',r2_score(y_train,i.predict(x_train)))
        print('Mean Squared Error is',mean_absolute_error(y_test,y_pred))
        print('Mean Squared Error is',mean_squared_error(y_test,y_pred))
        print('Root Mean Squared Error is',(mean_squared_error(y_test,y_pred,squared=False)))

KNeighborsRegressor()
R2 Score is 0.5723008556363665
R2 Score for train data 0.6651092826489714
Mean Squared Error is 1814.9600374356573
Mean Squared Error is 9041141.594010295
Root Mean Squared Error is 3006.8491139414186
SVR()
R2 Score is -0.03251945438190873
R2 Score for train data -0.02447186805496515
Mean Squared Error is 3627.1188577608436
Mean Squared Error is 21826451.393833652
Root Mean Squared Error is 4671.878786295044
DecisionTreeRegressor()
R2 Score is 0.6380395836239183
R2 Score for train data 0.7321334660000298
Mean Squared Error is 1741.7595521859014
Mean Squared Error is 7651489.181144442
Root Mean Squared Error is 2766.132531377418
```

```
from sklearn.model_selection import cross_val_score
```

```
for i in range(2,5):
    cv=cross_val_score(rfr,x,y,cv=i)
    print(rfr,cv.mean())
```

```
RandomForestRegressor() 0.6055338682850067
RandomForestRegressor() 0.6231738744144611
RandomForestRegressor() 0.6389578943109726
```

```
from sklearn.model_selection import RandomizedSearchCV
```

```
param_grid={'n_estimators':[10,30,50,70,100], 'max_depth':[None,1,2,3], 'max_features':['auto'],'sc
```

```
rfr=RandomForestRegressor()
```

```
rfr.fit(x_train,y_train)
```

```
▼ RandomForestRegressor
RandomForestRegressor()
```

```
gb=GradientBoostingRegressor()
```

```
gb_res=RandomizedSearchCV(estimator=gb,param_distributions=param_grid,cv=3,verbose=2,n_jobs=-1)
```

```
gb_res.fit(x_train,y_train)
```

```
Fitting 3 folds for each of 10 candidates, totalling 30 fits
```

```
► RandomizedSearchCV
► estimator: GradientBoostingRegressor
  ► GradientBoostingRegressor
```

Accuracy

```
rfr=RandomForestRegressor(n_estimators=10,max_features='sqrt',max_depth=None)
rfr.fit(x_train,y_train)
y_train_pred=rfr.predict(x_train)
y_test_pred=rfr.predict(x_test)
print("train accuracy",r2_score(y_train_pred,y_train))
print("test accuracy",r2_score(y_test_pred,y_test))
```

```
train accuracy 0.622418136124503
test accuracy 0.4663663616316054
```

```
knn=KNeighborsRegressor(n_neighbors=2,algorithm='auto',metric_params=None,n_jobs=-1)
knn.fit(x_train,y_train)
y_train_pred=knn.predict(x_train)
y_test_pred=knn.predict(x_test)
print("train accuracy",r2_score(y_train_pred,y_train))
print("test accuracy",r2_score(y_test_pred,y_test))
```

```
train accuracy 0.5082560160438218
test accuracy 0.28385693588182304
```

```
predicted_values=pd.DataFrame({'Actual':y_test,'Predicted':y_pred})
```

```
predicted_values
```

	Actual	Predicted	
6075	16655	19069.857143	
3544	4959	5496.333333	
9291	9187	8928.000000	
5032	3858	3657.230769	
2483	12898	12821.529412	
...	
9797	7408	12319.470588	
9871	4622	4903.714286	
10063	7452	7104.100000	
8803	7060	6244.185185	
8618	13731	11612.809524	

2137 rows × 2 columns

```
prices=rfr.predict(x_test)
```

```
price_list=pd.DataFrame({'Price':prices})
```

```
price_list
```

	Price	
0	19999.841747	
1	5604.342171	
2	8959.272976	
3	3654.219529	
4	13165.037332	
...	...	
2132	12589.262280	
2133	4764.224621	
2134	6873.395159	
2135	6253.708491	
2136	11819.839555	


2137 rows × 1 columns

```
import numpy as np
from matplotlib import pyplot as plt

ys = 200 + np.random.randn(100)
x = [x for x in range(len(ys))]

plt.plot(x, ys, '-')
plt.fill_between(x, ys, 195, where=(ys > 195), facecolor='g', alpha=0.6)

plt.title("Sample Visualization")
plt.show()
```



✓ 0s completed at 7:24 AM

